

Bootstrapping a Libre, Self-Hosting RISC-V Computer

Gabriel L. Somlo, Ph.D.

CERT / SEI
Carnegie Mellon University
Pittsburgh, PA 15213

Copyright 2021 Carnegie Mellon University.

This material is based upon work funded and supported by the Department of Defense under Contract No. FA8702-15-D-0002 with Carnegie Mellon University for the operation of the Software Engineering Institute, a federally funded research and development center.

The view, opinions, and/or findings contained in this material are those of the author(s) and should not be construed as an official Government position, policy, or decision, unless designated by other documentation.

NO WARRANTY. THIS CARNEGIE MELLON UNIVERSITY AND SOFTWARE ENGINEERING INSTITUTE MATERIAL IS FURNISHED ON AN "AS-IS" BASIS. CARNEGIE MELLON UNIVERSITY MAKES NO WARRANTIES OF ANY KIND, EITHER EXPRESSED OR IMPLIED, AS TO ANY MATTER INCLUDING, BUT NOT LIMITED TO, WARRANTY OF FITNESS FOR PURPOSE OR MERCHANTABILITY, EXCLUSIVITY, OR RESULTS OBTAINED FROM USE OF THE MATERIAL. CARNEGIE MELLON UNIVERSITY DOES NOT MAKE ANY WARRANTY OF ANY KIND WITH RESPECT TO FREEDOM FROM PATENT, TRADEMARK, OR COPYRIGHT INFRINGEMENT.

[DISTRIBUTION STATEMENT A] This material has been approved for public release and unlimited distribution. Please see Copyright notice for non-US Government use and distribution.

This material may be reproduced in its entirety, without modification, and freely distributed in written or electronic form without requesting formal permission. Permission is required for any other use. Requests for permission should be directed to the Software Engineering Institute at permission@sei.cmu.edu.

Carnegie Mellon® and CERT® are registered in the U.S. Patent and Trademark Office by Carnegie Mellon University.

DM21-0251

The Plan

- Build an FPGA based computer from Free / Libre software *and* gateway sources:
 - Linux, BusyBox
 - LiteX, RocketChip
- Using Free / Libre software *and* HDL toolchains:
 - gcc
 - yosys, trellis, nextpnr

You Can Follow Along

- On a pre-configured Fedora VM
 - <http://mirror.ini.cmu.edu/litex/litexdemo.f32.ova>
 - Built for VMWare (Fusion / Workstation), for convenience
 - Link availability *not* guaranteed beyond April 2021!
 - Login: *user*
 - Password: *tartans*
 - Pre-installed with toolchains, sources

Building the Development VM

- Kickstart a Fedora VM: [litedemo.f32.ks](#)
- Log in as *user* (password: *tartans*), and run:

```
# git clone -recursive https://github.com/riscv/riscv-gnu-toolchain
# pushd riscv-gnu-toolchain; configure -prefix=$HOME/RISCV -enable-multilib
# make newlib linux; popd; rm -rf riscv-gnu-toolchain
# echo 'export PATH=$PATH:$HOME/RISCV/bin' >> ~/.bashrc

# mkdir ~/LITEX; cd ~/LITEX
# git clone https://github.com/litex-hub/linux -b litex-rebase
# git clone https://github.com/litex-hub/linux-on-litex-rocket
# git clone https://github.com/riscv/riscv-pk
# wget https://raw.githubusercontent.com/enjoy-digital/litex/master/litex\_setup.py
# python3 ./litex_setup.py init install -user
# wget https://busybox.net/downloads/busybox-1.31.0.tar.bz2 -O - | tar xvj -
```

Let's start building!

- We'll talk *theory* while waiting for build to finish!
- Start with the bitstream (it takes the longest):

```
# cd ~/LITEX
# litex_boards/litex_boards/targets/lambdaconcept_ecp5.py --build \
  --cpu-type rocket --cpu-variant linuxd --sys-clk-freq 50e6 \
  --with-ethernet --with-sdcard
```

- Next, start building BusyBox:

```
# cp linux-on-litex-rocket/conf/busybox-*.config busybox*/.config
# (cd busybox*; make CROSS_COMPILE=riscv64-unknown-linux-gnu-)
```

Create the *initramfs.cpio* image

```
# mkdir initramfs; pushd initramfs; mkdir -p \  
    bin sbin lib etc dev home proc sys tmp mnt nfs root usr/bin usr/sbin usr/lib  
  
# cat > etc/inittab <<- "EOT"  
::sysinit:/bin/busybox mount -t proc proc /proc  
::sysinit:/bin/busybox mount -t tmpfs tmpfs /tmp  
::sysinit:/bin/busybox mount -t sysfs sysfs /sys  
::sysinit:/bin/busybox --install -s  
/dev/console::sysinit:-/bin/ash  
EOT  
  
# cp ../busybox*/busybox bin/; ln -s bin/busybox ./init  
  
# fakeroot <<- "EOT"  
mknod dev/null c 1 3; mknod dev/zero c 1 5; mknod dev/tty c 5 0  
mknod dev/console c 5 1; mknod dev/mmcblk0 b 179 0  
mknod dev/mmcblk0p1 b 179 1; mknod dev/mmcblk0p2 b 179 2  
find . | cpio -H newc -o > ../initramfs.cpio  
EOT  
# popd
```

Start building Linux

- Embed *initramfs.cpio* into the kernel image:

```
# cp initramfs.cpio linux/; pushd linux  
  
# make ARCH=riscv CROSS_COMPILE=riscv64-unknown-linux-gnu- \  
    litex_rocket_defconfig litex_rocket_initramfs.config  
  
# make ARCH=riscv CROSS_COMPILE=riscv64-unknown-linux-gnu- -j3  
  
# popd
```

- Let's discuss the *bigger picture* while we wait...

Self-Hosting (Compiler)

C compiler

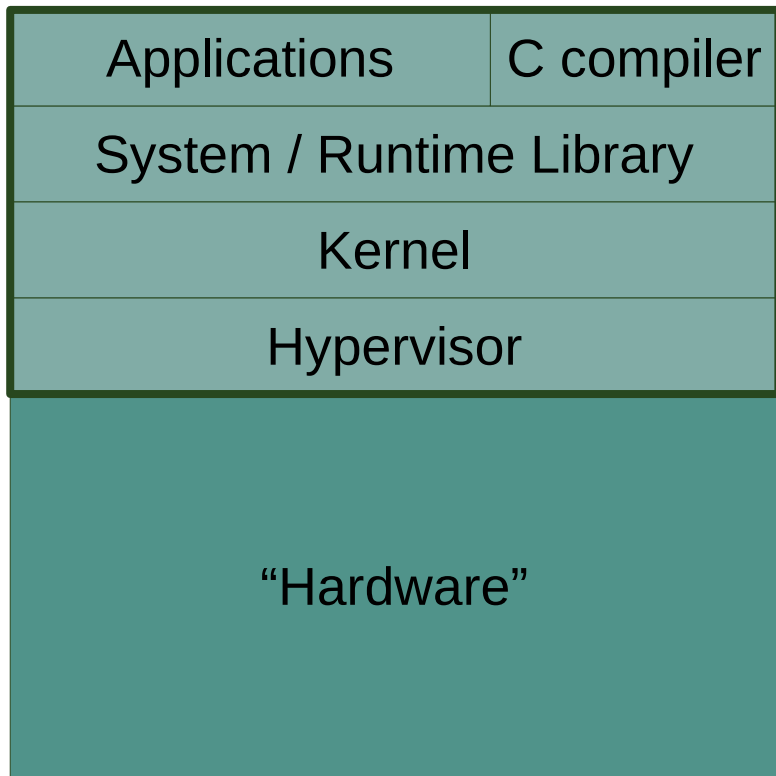
- Written in its own language
- Compiles its own sources
 - Subject to *Trusting Trust* attack (Ken Thompson)
 - Mitigated by *Diverse Double Compilation* (D. A. Wheeler)
- *Bootstrapping*

Self-Hosting Software Stack

| | |
|--------------------------|------------|
| Applications | C compiler |
| System / Runtime Library | |
| Kernel | |
| Hypervisor | |

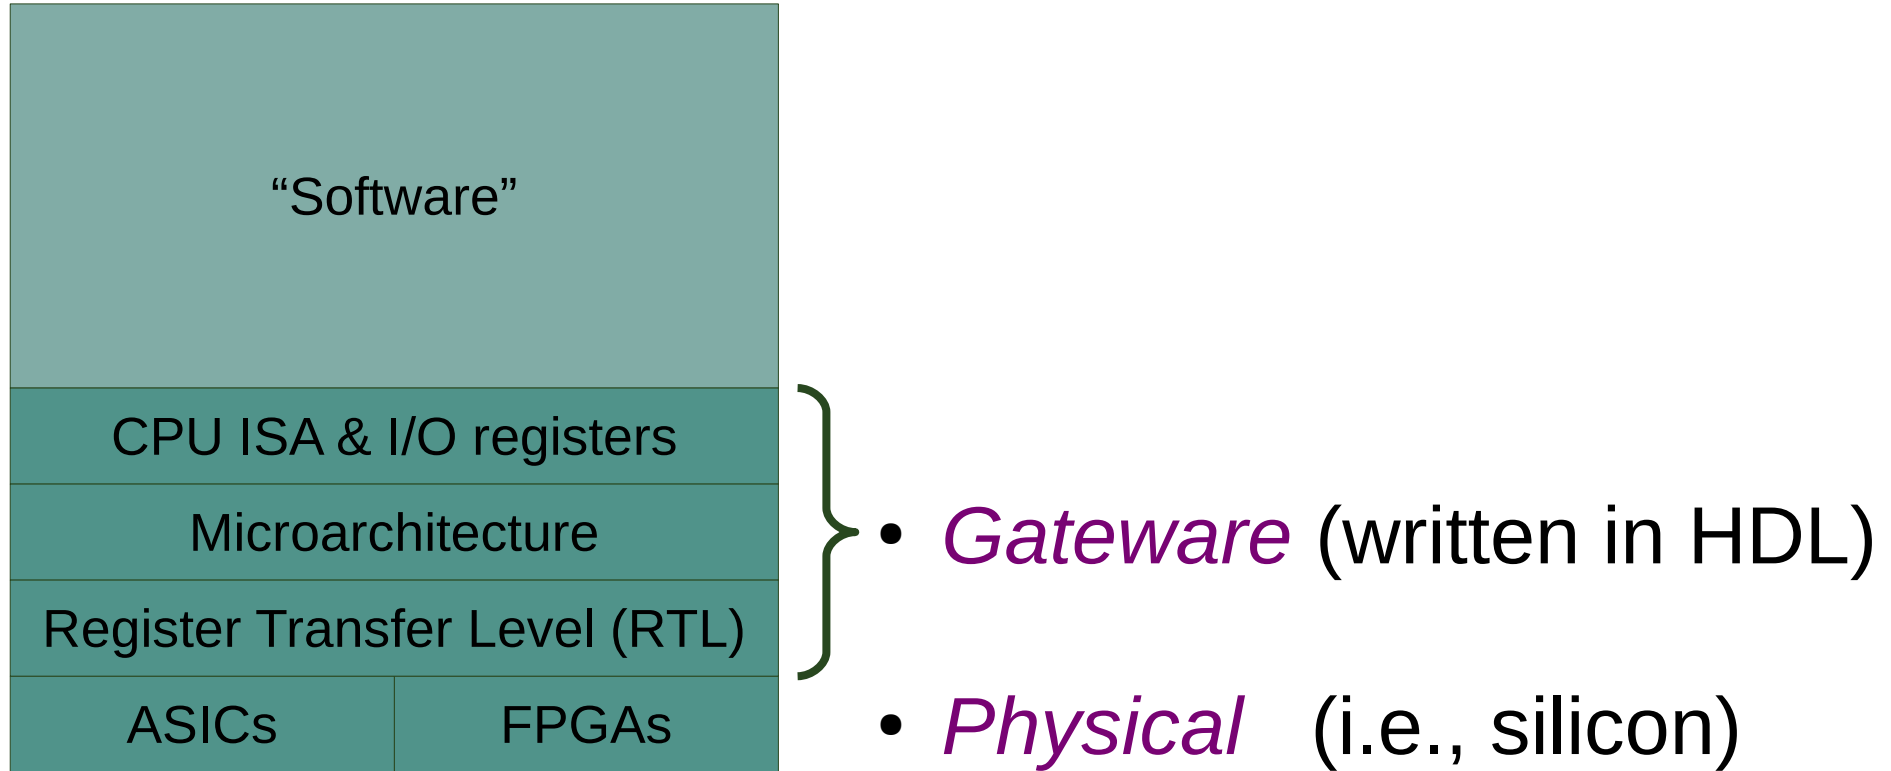
- Self-hosting compiler can build all software needed for its own execution
- Free / Libre sources for all components!

Self-Hosting Software Stack

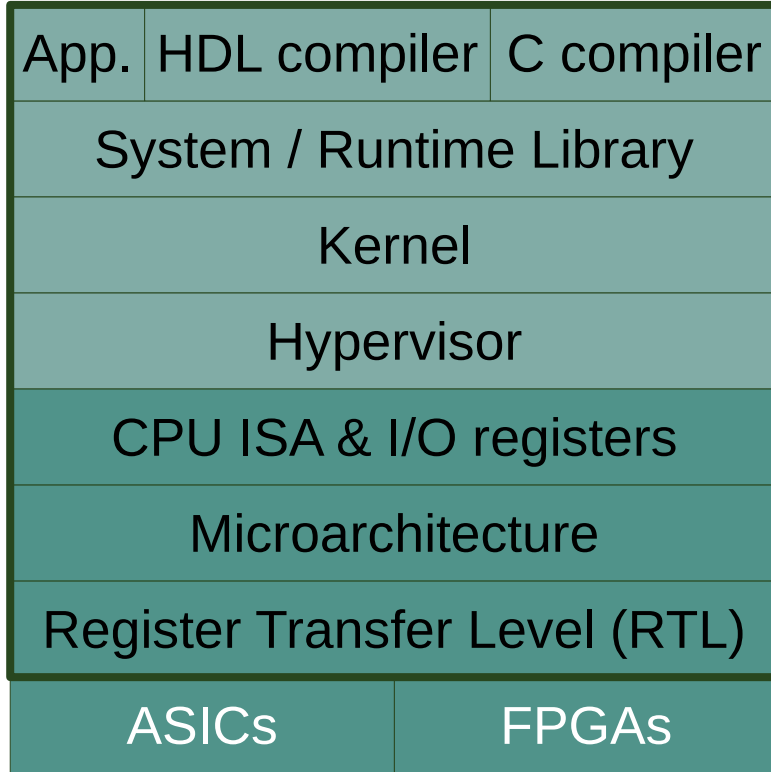


- Self-hosting compiler can build all software needed for its own execution
- Relies on (deployed on top of) *Hardware*

More Details re. *Hardware*



Self-Hosting Extended to Gateware



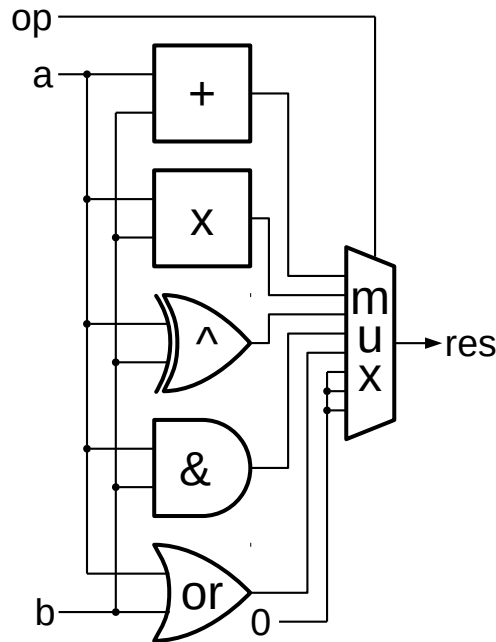
- C compiler → *Software*
- HDL compiler → *Gateware*
- Free / Libre sources for all components!
- *Physical* (silicon, ASICs or FPGAs) out of scope!

Gateway Compilation Stages

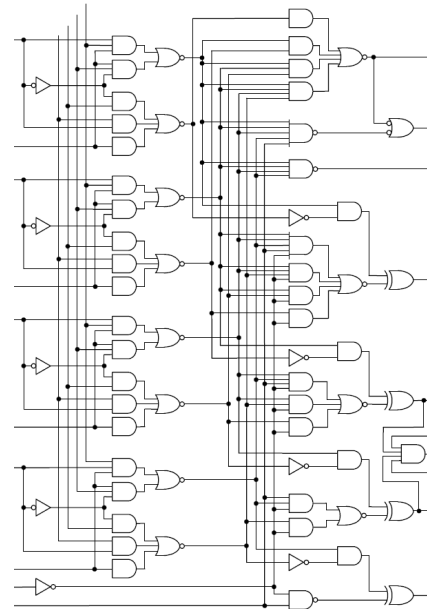
HDL Sources

```
module alu_mod (  
  // operator:  
  input  alu_op_t      op,  
  // operands:  
  input  logic [31:0] a, b,  
  // result:  
  output logic [31:0] res);  
  
  always_comb begin  
    unique case (op)  
      ALU_ADD: res = a + b;  
      ALU_MUL: res = a * b;  
      ALU_XOR: res = a ^ b;  
      ALU_AND: res = a & b;  
      ALU_OR : res = a | b;  
      default: res = 32'b0;  
    endcase  
  end  
endmodule: alu_mod
```

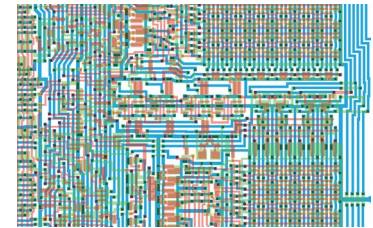
Elaboration



Synthesis,
Optimization



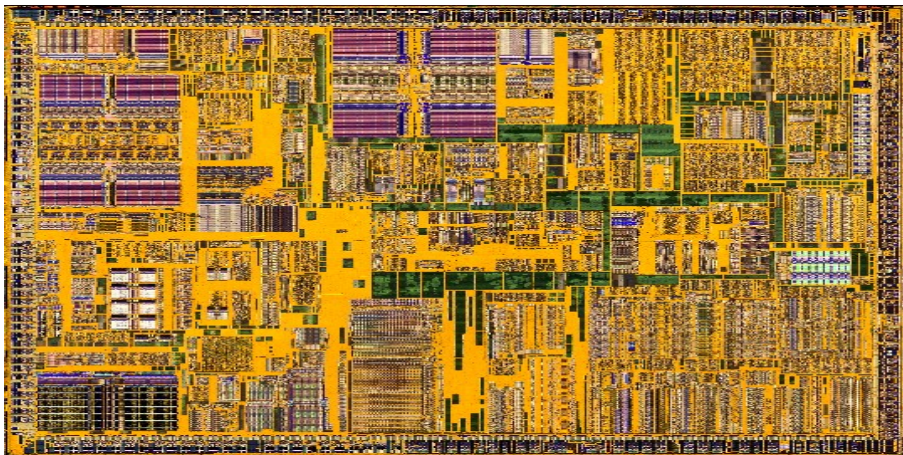
Tech. Mapping,
Place & Route



mask (ASIC)



ASICs vs. FPGAs



- Application Specific Integrated Circuit
- dedicated, optimized etched silicon
 - photolithographic masks
- *hard IP* cores



- Field Programmable Gate Array
- grid: programmable blocks, interconnect
 - bitstream
- *soft IP* cores

Hardware Attack Surface

- Fabrication (Malicious ASIC Foundry)
 - masks reverse engineered, modified to insert malicious behavior into ASIC
 - privilege escalation CPU backdoor ([A2 Trojan](#))
 - tamper with silicon [doping polarity](#) (e.g., to weaken hardware-based crypto)
 - problematic to test / verify after the fact!
 - mitigated by using FPGAs: hard to predict where to add *useful* Trojan silicon!
- Compilation ([Malicious HDL Toolchain](#))
 - generate *malicious* design from *clean* HDL sources
- Design Defects (accidental or intentional HDL bugs)
 - [Spectre](#), [Meltdown](#), etc.

Why *Self-Hosting* ?

- Freedom! Liberty! Independence! :)
 - From *black-box*, and/or *non-Libre* dependencies
- Trust a running *software + gateway stack* to the same extent as its *cumulative sources*
 - Gateway HDL sources
 - Software sources (including C and HDL compilers)

Bootstrap Software+Gateway Stack

- *Host* (x86_64/Linux):
 - Build clean C (cross-)compiler
 - Build clean HDL compiler (for both x86_64 and rv64gc)
 - Cross-compile target (rv64gc) software stack
 - Build gateway (FPGA bitstream) for target system
- *Target* (rv64gc/Linux):
 - Program FPGA board with gateway/bitstream
 - Boot into target software stack
 - Self-hosting from this point forward!
 - Natively rebuild gateway bitstream, software stack, from sources, as needed

LiteX + Rocket SoC Block Diagram

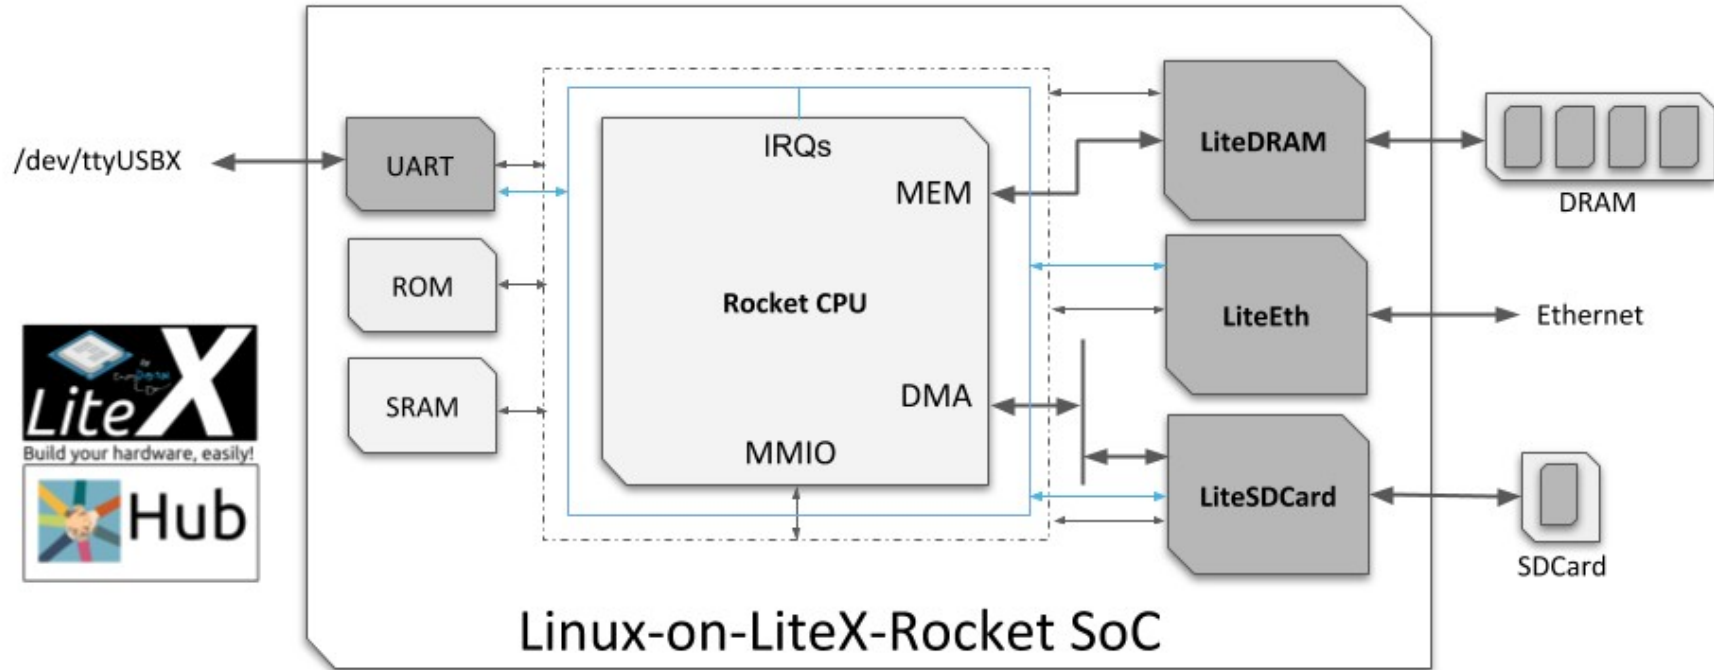


Image credit: [Florent Kermarrec](#)

Finish building *boot.bin* image

- Embed kernel into boot.bin (BBL):

```
# mkdir riscv-pk/build; pushd riscv-pk/build  
  
# ../configure host=riscv64-unknown-linux-gnu --enable-logo --with-arch=rv64imac  
    --with-payload=../../linux/vmlinux  
    --with-dts=../../linux-on-litex-rocket/conf/ecpix5.dts  
  
# make bbl  
  
# riscv64-unknown-linux-gnu-objcopy -O binary bbl ../../boot.bin  
  
# popd
```

- Make *boot.bin* available via TFTP, or copy to 1st MSDOS / FAT16 primary partition of SDCard

Boot Linux on LiteX+Rocket SoC

- Connect board via USB, and start a console:

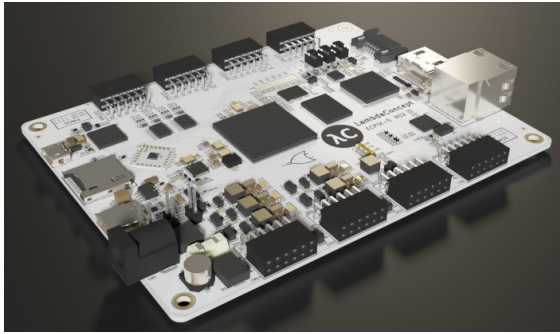
```
# screen dev/ttyUSB1 115200
```

- Program the board with the compiled bitstream:

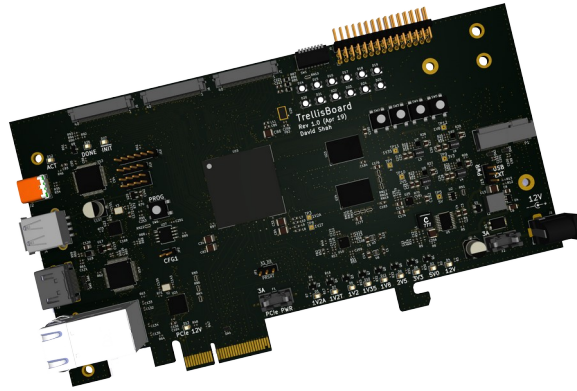
```
# openocd -f litex-boards/litex_boards/prog/openocd_ecp5.cfg \  
-c 'transport select jtag; init;  
svf build/lambdaconcept_ecp5/gateway/lambdaconcept_ecp5.svf; exit'
```

- LiteX loads *boot.bin*, Linux boots into BusyBox

Try it on your own FPGA board!



ECPIX-5



trellisboard



ecp5-5g-versa

Demo, then Q&A

- For up-to-date build steps, see:
<https://github.com/litex-hub/linux-on-litex-rocket>
- Thank You!

Backup Material

Compilers, Trusting Trust, and DDC

- Ken Thompson's **self-propagating C compiler hack**
 - malicious compiler inserts Trojan during compilation of *victim program*
 - clean sources → malicious binary (incl. *compiler's own sources*!)
 - compiler source hack *no longer needed* beyond 1st iteration!
- David A. Wheeler's mitigation: **Diverse Double Compilation (DDC)**
 - suspect compiler A: sources S_A , binary B_A
 - trusted compiler T: binary B_T

$$S_A \rightarrow B_A \rightarrow X$$

- X and Y are *functionally identical*, but *different binaries*

$$S_A \rightarrow X \rightarrow X_1$$

- X_1 and Y_1 must be *identical binaries* (output of two *functionally identical* compilers)!

$$S_A \rightarrow B_T \rightarrow Y$$

$$S_A \rightarrow Y \rightarrow Y_1$$